

Saketh Sree Ram

Bangalore, India | +91-93467 48201 | reachout.sreeram@gmail.com | linkedin.com/in/sakethsreeram7 | github.com/reachout-sreeram

Summary

AI/ML Engineer, 2+ years building production LLM and agentic systems at scale. Shipped a fine-tuned DeBERTa (**99.8% F1, 1000+ enterprise customers**) and an agentic assistant that cut developer workload **75%**; **Winner, AWS Gen AI Hackathon 2025**. Expertise across RAG, MCP-based agent orchestration, LLM evaluation/observability, and end-to-end LLMops on AWS.

Skills

ML & LLMs: PyTorch, CatBoost, DeBERTa, Hugging Face, Fine-tuning, Scikit-learn

GenAI & Agents: Retrieval-Augmented Generation (RAG), Agentic Systems, Multi-Agent Orchestration, Human-in-the-Loop, LangChain, LangGraph, CrewAI, AWS Bedrock, Model Context Protocol (MCP), ChromaDB, Knowledge Graphs

LLMOps & Evaluation: LLM-as-Judge Evaluation, Golden Datasets, Guardrails, Hallucination Detection, Tracing & Observability, Cost Attribution, Context Engineering

Cloud & MLOps: AWS (SageMaker, Lambda, ECR), Docker, Git, CI/CD, GPU Inference

Languages & Data: Python, Node.js, TypeScript, REST APIs, PostgreSQL, Pandas, NumPy, Data Pipelines

Work Experience

Aquera

AI/ML Engineer

Apr 2024 – Present

Bangalore, India

- **Agentic AI Assistant (AWS Bedrock):** Built a Claude-powered agentic loop with MCP connectors, human-in-the-loop gating, and multi-agent orchestration (LangGraph, CrewAI-style task decomposition) for HR-to-AD script generation; context compression for nested API schemas. Cut script-writing **75%** across **100+ developers**.
- **DeBERTa Fine-tuning & Evaluation:** Fine-tuned DeBERTa-v3 (40K HR records) for service-account classification; built an LLM-assisted labeling pipeline and golden-dataset eval harness with per-slice F1 tracking. **99.8% F1** on SageMaker, serving **1000+ enterprise customers**.
- **IAM Rule-Mining Engine:** Built a CatBoost + RuleFit pipeline for auditable role-assignment rules, modeled as a knowledge graph. **85%+ accuracy** on 350 live datasets; cut onboarding from **4–8 weeks to a 1–2 hour call** (SageMaker, ECR, GPU inference).
- **GenBI Platform:** Built a RAG-based NL→SQL engine with LLM-generated dashboards for self-serve analytics, eliminating engineering dependency for routine queries.
- **MLOps & LLMOps:** Owned end-to-end deployment across all ML systems (Docker, CI/CD, AWS Lambda, SageMaker, ECR); LLM observability via call tracing, cost attribution, and guardrails over PostgreSQL dataflows.

Projects

VS Code Marketplace Extensions | *TypeScript, VS Code Extension API* | 2026

- **Variable Explorer:** Published a VS Code extension with a hybrid TypeScript–Python IPC architecture: a gateway process bridges the extension host to a live IPython kernel over ZeroMQ via newline-delimited JSON on stdin/stdout; stateful webviews with paginated DataFrame views.
- **Focus Timer (Agent Memory):** Designed a local `.focus/` persistence layer (JSON/Markdown) driven by extension lifecycle and file-system/git watchers — a recoverable workspace memory shared between developers and AI coding agents.

AI Job Search Agent | *Claude Agent SDK, FastAPI, Next.js, LangChain, ChromaDB* | 2026

- Designed an asynchronous multi-agent system: FastAPI backend with background workers on a job queue, a ChromaDB vector layer for profile/JD matching and caching, and a Next.js Kanban UI for human-in-the-loop review. Agents handle discovery, RAG form-filling, and email generation; cut per-application time **90%**.

Code Compass | *VS Code Extension, AWS Bedrock, MCP* | 2025

- Built an MCP integration layer (team of 5) unifying GitHub, Slack, and Jira as live context sources behind a VS Code extension; orchestrated retrieval across all three for component-level code walkthroughs. Cut onboarding **75%**.

Education

New Horizon College of Engineering

B.E., Artificial Intelligence and Machine Learning

Jun 2024

GPA: **9.57 / 10**

Achievements

- **Winner, AWS Gen AI Hackathon 2025** — For engineering quality and novelty of MCP-based agentic architecture (team of 5)